

Rectangle Filters and AdaBoost

CSE 6367 – Computer Vision
Vassilis Athitsos
University of Texas at Arlington

A Generalized View of Classifiers

- We have studied two face detection methods:
 - Normalized correlation.
 - PCA.
- Each approach exhaustively evaluates all image subwindows.
- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.

Features and Classifiers

- Our goal, in the next slides, is to get a better understanding of:
 - What is a feature?
 - What is a classifier?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?
- Two possible answers:
 - Each pixel value is a feature.
 - The feature is the result of normalized correlation with the template.

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?
 - The result of normalized correlation with the template.
 - Arguably, the score is the feature itself.

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?
 - In `find_template`, faces are the top N scores.
 - N is an argument to the `find_template` function.
 - An alternative, is to check if $\text{score} > \text{threshold}$.
 - Then we must choose a threshold instead of N.

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?
 - The result of the `pca_projection` function.
 - In other words, the K numbers that describe the projection of the subwindow on the space defined by the top K eigenfaces.

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?
 - The result of the `pca_score` function.
 - The sum of differences between:
 - the original subwindow W , and
 - `pca_backprojection(pca_projection(W))`.

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?
 - In `find_faces`, faces are the top N scores.
 - N is an argument to the `find_faces` function.
 - An alternative, is to check if $\text{score} > \text{threshold}$.
 - Then we must choose a threshold instead of N.

Defining a Classifier

- Choose features.
- Choose a scoring function.
- Choose a decision process.
 - The last part is usually straightforward.
 - We pick the top N scores, or we apply a threshold.
- Therefore, the two crucial components are:
 - Choosing features.
 - Choosing a scoring function.

What is a Feature?

- Any information extracted from an image.
- The number of possible features we can define is enormous.
- Any function F we can define that takes in an image as an argument and produces one or more numbers as output defines a feature.
 - Correlation with a template defines a feature.
 - Projecting to PCA space defines features.
 - More examples?

What is a Feature?

- Any information extracted from an image.
- The number of possible features we can define is enormous.
- Any function F we can define that takes in an image as an argument and produces one or more numbers as output defines a feature.
 - Correlation with a template defines a feature.
 - Projecting to PCA space defines features.
 - Average intensity.
 - Std of values in window.

Boosted Rectangle Filters

- Boosting is a machine learning method.
- Rectangle filters provide us with an easy way to define features.
- Boosted rectangle filters is an extremely popular computer vision method that answers the basic two questions in classifier design:
 - What features do we use? Rectangle filters.
 - What scoring function do we use? The result of boosting.

Boosted Rectangle Filters

- Boosted rectangle filters is an extremely popular computer vision method that answers the basic two questions in classifier design:
 - What features do we use? Rectangle filters.
 - What scoring function do we use? The result of boosting.
- The popularity of this method is due to two factors:
 - Can be applied in lots of problems.
 - Works well in lots of problems.

What is a Rectangle Filter?

- $\begin{bmatrix} 1 & 1 & 1 & 1 & -1 & -1 & -1 & -1 \\ 1 & 1 & 1 & 1 & -1 & -1 & -1 & -1 \\ 1 & 1 & 1 & 1 & -1 & -1 & -1 & -1 \\ 1 & 1 & 1 & 1 & -1 & -1 & -1 & -1 \end{bmatrix}$

- $\begin{bmatrix} 1 & 1 & 1 & -2 & -2 & -2 & 1 & 1 & 1 \\ 1 & 1 & 1 & -2 & -2 & -2 & 1 & 1 & 1 \\ 1 & 1 & 1 & -2 & -2 & -2 & 1 & 1 & 1 \\ 1 & 1 & 1 & -2 & -2 & -2 & 1 & 1 & 1 \\ 1 & 1 & 1 & -2 & -2 & -2 & 1 & 1 & 1 \\ 1 & 1 & 1 & -2 & -2 & -2 & 1 & 1 & 1 \end{bmatrix}$

- $\begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 \end{bmatrix}$

- $\begin{bmatrix} 1 & 1 & 1 & -1 & -1 & -1 \\ 1 & 1 & 1 & -1 & -1 & -1 \\ 1 & 1 & 1 & -1 & -1 & -1 \\ 1 & 1 & 1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \\ -1 & -1 & -1 & 1 & 1 & 1 \end{bmatrix}$

Five Different Basic Shapes

- Type 1:
 - Two areas, horizontal.

```
[1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1]
```



white: value = 1
black: value = -1

- To define:
 - Specify rectangle size:
 - number of rows
 - number of cols.

Five Different Basic Shapes

- Type 1:
 - Two areas, horizontal.

```
[1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1]
```



white: value = 1
black: value = -1


- To define:
 - Specify rectangle size:
 - number of rows
 - number of cols.

Have we ever used anything similar?

Five Different Basic Shapes

- Type 1:
 - Two areas, horizontal.

```
[1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1
 1 1 1 1 -1 -1 -1 -1]
```



white: value = 1
black: value = -1


- To define:
 - Specify rectangle size:
 - number of rows
 - number of cols.

Have we ever used anything similar?
 $dx = [-1 \ 0 \ -1]$ is a centered version of `rectangle_filter1(1, 1)`.

Five Different Basic Shapes

- Type 2:
 - Two areas, vertical.

```
[ 1  1  1  1  1  1
  1  1  1  1  1  1
  1  1  1  1  1  1
 -1 -1 -1 -1 -1 -1
 -1 -1 -1 -1 -1 -1
 -1 -1 -1 -1 -1 -1]
```



To define:

- Specify rectangle size:
 - number of rows
 - number of cols.

Five Different Basic Shapes

- Type 2:
 - Two areas, vertical.

```
[ 1  1  1  1  1  1
  1  1  1  1  1  1
  1  1  1  1  1  1
 -1 -1 -1 -1 -1 -1
 -1 -1 -1 -1 -1 -1
 -1 -1 -1 -1 -1 -1]
```

To define:

- Specify rectangle size:
 - number of rows
 - number of cols.



Have we ever used anything similar?


Five Different Basic Shapes

- Type 2:
 - Two areas, vertical.

```
[ 1  1  1  1  1  1
  1  1  1  1  1  1
  1  1  1  1  1  1
 -1 -1 -1 -1 -1 -1
 -1 -1 -1 -1 -1 -1
 -1 -1 -1 -1 -1 -1]
```

To define:

- Specify rectangle size:
 - number of rows
 - number of cols.



Have we ever used anything similar?

$dy = [-1 \ 0 \ -1]^T$ is a centered version of `rectangle_filter2(1, 1)`.

Effect of Rectangle Size

- Larger rectangles \rightarrow ?

Effect of Rectangle Size

- Larger rectangles → emphasis on larger scale features.
- Smaller rectangles → emphasis on smaller scale features.

Five Different Basic Shapes

- Type 3:
 - Three areas, horizontal.

```
[1  1  -2  -2  1  1
 1  1  -2  -2  1  1
 1  1  -2  -2  1  1
 1  1  -2  -2  1  1]
```



white: value = 1
black: value = -2

- To define:
 - Specify rectangle size:
 - number of rows
 - number of cols.

Five Different Basic Shapes

- Type 4:
 - Three areas, vertical.

```
[1  1  1  1
 1  1  1  1
-2 -2 -2 -2
-2 -2 -2 -2
 1  1  1  1
 1  1  1  1]
```




white: value = 1
black: value = -2

- To define:
 - Specify rectangle size:
 - number of rows
 - number of cols.

Five Different Basic Shapes

- Type 5:
 - Four areas, diagonal.

```
[1  1  1 -1 -1 -1
 1  1  1 -1 -1 -1
-1 -1 -1  1  1  1
-1 -1 -1  1  1  1]
```



- To define:
 - Specify rectangle size:
 - number of rows
 - number of cols.

white: value = 1
black: value = -1

Advantages of Rectangle Filters

- Easy to define and implement.
- Lots of them.
 - Intuitively, for many patterns that we want to detect/recognize, we can find some rectangle filters that are useful.
- Fast to compute.
 - How?

Advantages of Rectangle Filters

- Easy to define and implement.
- Lots of them.
 - Intuitively, for many patterns that we want to detect/recognize, we can find some rectangle filters that are useful.
- Fast to compute.
 - How?
 - Using integral images (defined in the next slides).

Integral Image

- Input: grayscale image A.
 - $A(i, j)$ is the intensity value at pixel (i, j) .
- Integral image B:
 - $B(i, j) = \text{sum}(\text{sum}(A(1:i, 1:j)))$.
- $F = \text{rectangle_filter1}(50, 40);$
- How can I compute the response on subwindow $A(101:150, 201:280)$?

Integral Image

- Input: grayscale image A.
 - $A(i, j)$ is the intensity value at pixel (i, j) .
- Integral image B.
 - $B(i, j) = \text{sum}(\text{sum}(A(1:i, 1:j)))$.
- $F = \text{rectangle_filter1}(50, 40);$
- How can I compute the response on subwindow $A(101:150, 201:280)$?
 - First approach:
 $\text{sum}(\text{sum}(A(101:150, 201:280) .* F))$
 - How many operations does that take? $O(50*80)$

Integral Image

- Input: grayscale image A.
- Integral image: $B(i, j) = \text{sum}(\text{sum}(A(1:i, 1:j)))$.
- $F = \text{rectangle_filter1}(50, 40);$
- How can I compute the response on subwindow $A(101:150, 201:280)$?
 - $\text{sum}(\text{sum}(A(101:150, 201:240))) - \text{sum}(\text{sum}(A(101:150, 241:280)))$.
 - where is F in the sum?

Integral Image


- Input: grayscale image A.
- Integral image: $B(i, j) = \text{sum}(\text{sum}(A(1:i, 1:j)))$.
- $F = \text{rectangle_filter1}(50, 40);$
- How can I compute the response on subwindow $A(101:150, 201:280)$?
 - $\text{sum}(\text{sum}(A(101:150, 201:240))) - \text{sum}(\text{sum}(A(101:150, 241:280)))$.
 - where is F in the sum?
 - F is encoded in specifying the submatrices that we sum.

Integral Image

- Integral image: $B(i, j) = \text{sum}(\text{sum}(A(1:i, 1:j)))$.
- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$
fast using integral image:


Integral Image

- Integral image: $B(i, j) = \text{sum}(\text{sum}(A(1:i, 1:j)))$.
- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$
fast using integral image:
 - $B(150,240)+B(100,200)-B(150,200)-B(100,240)$




Integral Image

- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$
fast using integral image:
– $B(150,240)+B(100,200)-B(150,200)-B(100,240)$




Integral Image

- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$ fast using integral image:
 - $B(150,240)+B(100,200)-B(150,200)-B(100,240)$
 - $\text{sum}(\text{Area1})+\text{sum}(\text{Area2})-\text{sum}(\text{Area3})-\text{sum}(\text{Area4})$




Integral Image

- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$ fast using integral image:
 - $B(150,240)+B(100,200)-B(150,200)-B(100,240)$
 - $\text{sum}(\text{Area1})+\text{sum}(\text{Area2})-\text{sum}(\text{Area3})-\text{sum}(\text{Area4})$




Integral Image

- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$ fast using integral image:
 - $B(150,240)+B(100,200)-B(150,200)-B(100,240)$
 - $\text{sum}(\text{Area1})+\text{sum}(\text{Area2})-\text{sum}(\text{Area3})-\text{sum}(\text{Area4})$



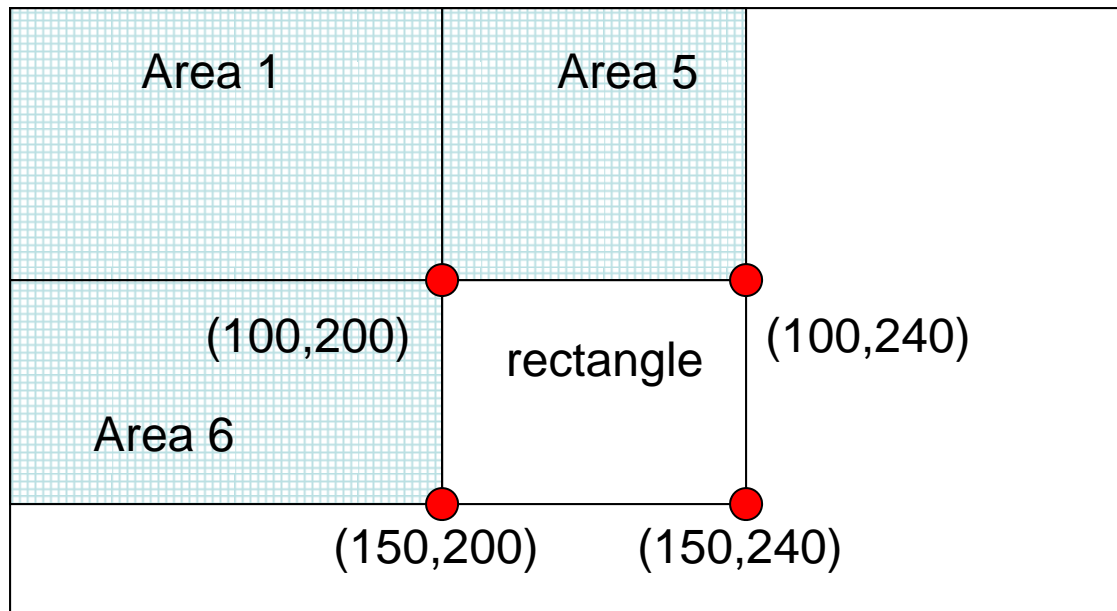
Integral Image

- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$ fast using integral image:
 - $B(150,240)+B(100,200)-B(150,200)-B(100,240)$
 - $\text{sum}(\text{Area1})+\text{sum}(\text{Area2})-\text{sum}(\text{Area3})-\text{sum}(\text{Area4})$



Integral Image

- Compute $\text{sum}(\text{sum}(A(101:150, 201:240)))$ fast using integral image:
 - $B(150,240)+B(100,200)-B(150,200)-B(100,240)$
 - $\text{sum}(\text{Area1})+\text{sum}(\text{Area2})-\text{sum}(\text{Area3})-\text{sum}(\text{Area4})$



- Area1: added twice, subtracted twice.
- Area5: added once, subtracted once.
- Area6: added once, subtracted once.
- Rectangle: added once.

Rectangle Filters for Face Detection

- Key questions in using rectangle filters for face detection:
 - Or for any other classification task.
- How do we use a rectangle filter to extract information?
- What rectangle filter-based information is the most important?
- How do we combine information from different rectangle filters?

From Rectangle Filter to Classifier

- We want to define a classifier, that “guesses” if a specific image window is a face or not.
 - Remember, all face detector operate window-by-window.
- Convention:
 - Classifier output +1 means “face”.
 - Classifier output -1 means “not a face”.

What is a Classifier?

- What is the input of a classifier?
- What is the output of a classifier?

What is a Classifier?

- What is the input of a classifier?
 - An entire image, or
 - An image subwindow, or
 - Any type of pattern (audio, biological, ...).
- What is the output of a classifier?

What is a Classifier?

- What is the input of a classifier?
 - An entire image, or
 - An image subwindow, or
 - Any type of pattern (audio, biological, ...).
- What is the output of a classifier?
 - A *class label*.
 - We have a *finite* number of classes.

What is a Classifier?

- Focusing on face detection:
- A classifier takes as input an image *subwindow*.
- A classifier outputs +1 (for “face”) or -1 (for “non-face”).
- What is a classifier?

What is a Classifier?

- Focusing on face detection:
- A classifier takes as input an image *subwindow*.
- A classifier outputs +1 (for “face”) or -1 (for “non-face”).
- What is a classifier?
 - Any function that takes the specified input and produces the specified output.
- VERY IMPORTANT: “classifier” is not the same as “accurate classifier.”

Classifier Based on a Rectangle Filter

- How can we define a classifier using a single rectangle filter?
- Goal:
 - We want something very fast.
 - It does not have to always be very accurate.
 - We will try LOTS of such classifiers, we are happy as long as SOME of them are very accurate.

Classifier Based on a Rectangle Filter

- How can we define a classifier using a single rectangle filter?
- Answer: choose response on a specific location.
- Classifier is specified by:
 - Type
 - Rectangle size
 - *Window offset.*
 - *Threshold.*


Classifier Based on a Rectangle Filter

- How can we define a classifier using a single rectangle filter?
- Answer: choose response on a specific location.
- Classifier is specified by:
 - Type
 - Rectangle size
 - *Window offset.*
 - *Threshold.*




Classifier Based on a Rectangle Filter

- How can we define a classifier using a single rectangle filter?
- Answer: choose response on a specific location.
- Classifier is specified by:
 - Type
 - Rectangle size
 - *Window offset* (row, col).
 - *Threshold*.




Classifier Based on a Rectangle Filter

- Classifier is specified by:
 - Type
 - Rectangle size
 - *Window offset* (row, col).
 - *Threshold*.
- How many classifiers do we get from a single rectangle filter?



Classifier Based on a Rectangle Filter

- Classifier is specified by:
 - Type
 - Rectangle size
 - *Window offset* (row, col).
 - *Threshold*.
- How many classifiers do we get from a single rectangle filter?
 - upper bound:
 $\# \text{ pixels} * \# \text{ thresholds}$



Useful Functions

- `generate_classifier1`
- `generate_classifier2`
- `generate_classifier3`
- `generate_classifier4`
- `generate_classifier`
- `eval_weak_classifier`

How Many Classifiers?

- Upper bound:
 - 31x25 offsets.
 - 5 types
 - 31x25 rectangle sizes.
 - Lots of different possible thresholds.
- Too many to evaluate.
- We can sample.
 - Many classifiers are similar to each other.
 - Approach: pick some (a few thousands).
 - Function: `generate_classifier`

Terminology

- For the purposes of this lecture, we will introduce the term *soft classifier*.
- A soft classifier is defined by specifying:
 - A rectangle filter.
 - (i.e., type of filter and size of rectangle).
 - An offset.
- A weak classifier is defined by specifying:
 - A soft classifier.
 - (i.e., rectangle filter and offset).
 - A threshold.

Precomputing Responses

- Suppose we have a training set of faces and non-faces.
- Suppose we have picked 1000 soft classifiers.
- For each training example, we can precompute the 1000 responses, and save them.
- This way, we know for each soft classifier and each example what the response is.

Precomputing Responses

- Suppose we have a training set of faces and non-faces.
- Suppose we have picked 1000 soft classifiers.
- For each training example, we can precompute the 1000 responses, and save them.
- This way, we know for each soft classifier and each example what the response is.
- How many classifiers does that correspond to?

Precomputing Responses

- Suppose we have a training set of faces and non-faces.
- Suppose we have picked 1000 soft classifiers.
- For each training example, we can precompute the 1000 responses, and save them.
- This way, we know for each soft classifier and each example what the response is.
- How many classifiers does that correspond to? upper bound: $1000 * \# \text{ thresholds}$.

How Useful is a Classifier?

- First version: how useful is a soft classifier by itself?
- Measure the error.
 - Involves finding an optimal threshold.
 - Also involves deciding if faces tend to be above or below the threshold.

```

function [error, best_threshold, best_alpha] = ...
    weighted_error(responses, labels, weights, classifier)

classifier_responses = [responses(classifier,:)];
minimum = min(classifier_responses);
maximum = max(classifier_responses);
step = (maximum - minimum) / 50;
best_error = 2;
for threshold = minimum:step:maximum
    thresholded = double(classifier_responses > threshold);
    thresholded(thresholded == 0) = -1;
    error1 = sum(weights .* (labels ~= thresholded));
    error = min(error1, 1 - error1);
    if (error < best_error)
        best_error = error;
        best_threshold = threshold;
        best_direction = 1;
        if (error1 > (1 - error1))
            best_direction = -1;
        end
    end
end
best_alpha = best_direction * 0.5 * log( (1 - error) / error);
if (best_error == 0)
    best_alpha = 1;
end

```

```
function [index, error, threshold, alpha] = ...  
    find_best_classifier(responses, labels, weights)  
  
classifier_number = size(responses, 1);  
example_number = size(responses, 2);  
  
% find best classifier  
best_error = 2;  
  
for classifier = 1:classifier_number  
    [error threshold alpha] = weighted_error(responses, labels, ...  
                                              weights, classifier);  
  
    if (error < best_error)  
        best_error = error;  
        best_threshold = threshold;  
        best_classifier = classifier;  
        best_alpha = alpha;  
    end  
end  
  
index = best_classifier;  
error = best_error;  
threshold = best_threshold;  
alpha = best_alpha;
```

How to Combine Many Classifiers

- 3% error is not great for a face detector.
- Why?

How to Combine Many Classifiers

- 3% error is not great for a face detector.
- Why?
 - Because the number of windows is huge:
 - $\text{\#pixels} * \text{\#scales}$, = hundreds of thousands or more.
 - We do not want 1000 or more false positives.
- We need to combine information from many classifiers.
 - How?

Boosting (AdaBoost)

- Weak classifier:
 - Better than a random guess.
 - What does that mean?

Boosting (AdaBoost)

- We have a binary problem.
 - Face or no face.
 - George or not George.
 - Mary or not Mary.
- Weak classifier:
 - Better than a random guess.
 - What does that mean?
 - For a binary problem, less than 0.5 error rate.
- Strong classifier:
 - A “good” classifier (definition intentionally fuzzy).

Boosted Strong Classifier

- Linear combination of weak classifiers.
- Given lots of classifiers:
 - Choose d of them.
 - Choose a weight α_d for each of them.
- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots$
 $+ \alpha_d * h_d(\text{pattern}).$
- Goal: strong classifier H should be significantly more accurate than each weak classifier h_1, h_2, \dots

Detour: Optimization

- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots$
 $+ \alpha_d * h_d(\text{pattern}).$
- How can we cast this as an optimization problem?

Detour: Optimization

- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots$
 $+ \alpha_d * h_d(\text{pattern}).$
- How can we cast this as an optimization problem?
 - Identify free parameters:

Detour: Optimization

- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots$
 $+ \alpha_d * h_d(\text{pattern}).$
- How can we cast this as an optimization problem?
 - Identify free parameters:
 - alphas (including zero weights for classifiers that are not selected).
 - Thresholds for the soft classifiers that were selected.

Detour: Optimization

- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots$
 $+ \alpha_d * h_d(\text{pattern}).$
- How can we cast this as an optimization problem?
 - Identify free parameters.
 - Define optimization criterion:

Detour: Optimization

- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots + \alpha_d * h_d(\text{pattern}).$
- How can we cast this as an optimization problem?
 - Identify free parameters.
 - Define optimization criterion:
 - $F(\text{parameters}) = \text{error rate attained on training set using those parameters.}$
 - We want to find optimal parameters for minimizing F .

Detour: Optimization

- Optimization:
 - Can be global or local.
 - Lots of times when people say “optimal” they really mean “locally optimal”.
- Given $F(\text{parameters})$ = error rate attained on training set using those parameters:
 - How can we do optimization?

Detour: Optimization

- Optimization:
 - Can be global or local.
 - Lots of times when people say “optimal” they really mean “locally optimal”.
- Given $F(\text{parameters}) = \text{error rate attained on training set using those parameters}$:
 - How can we do optimization?
 - Try many different combinations.
 - Gradient descent.
 - Greedy: choose weak classifiers one by one.
 - First, choose best weak classifier.
 - Then, at each iteration, choose next weak classifier and weight, to be the one giving the best results when combined with the previously picked classifiers and weights.

Boosted Strong Classifier

- Linear combination of weak classifiers.
- Given lots of classifiers:
 - Choose d of them.
 - Choose a weight α_d for each of them.
- Final result:
 - $H(\text{pattern}) = \alpha_1 * h_1(\text{pattern}) + \dots$
 $+ \alpha_d * h_d(\text{pattern}).$
- Goal: strong classifier H should be significantly more accurate than each weak classifier h_1, h_2, \dots

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0:

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0:
 - $\log(1/0) = \log(\text{inf}) = \text{inf}$.
 - depending on direction, $\alpha = +\text{inf}$ or $-\text{inf}$

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0.5:

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0.5:
 - $\log(0.5/0.5) = \log(1) = ?$

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0.5:
 - $\log(0.5/0.5) = \log(1) = 0$.
 - Capturing the fact that the weak classifier is useless.

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0.9:

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0.9:
 - $\log(0.1/0.9) = \log(.111) = ?$

Choosing the First Weak Classifier

- First task: choose α_1, h_1 .
 - Use function `find_best_classifier`.
 - Choose an α_1 for that classifier.
 - Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Interpreting alpha formula:
 - When error is 0.9:
 - $\log(0.1/0.9) = \log(.111) = -2.197$
 - Capturing the fact that we should do the opposite of what the classifier says.

Updating Training Weights

- So far we have chosen *alpha_1* and *h1*.
- Next step: update training weights.
- Initially, training weights are uniform:
 - each weight is $1/(\text{number of training examples})$.

Updating Training Weights

- So far we have chosen *alpha_1* and *h1*.
- Next step: update training weights.
- Initially, training weights are uniform:
 - each weight is $1/(\text{number of training examples})$.
- **VERY IMPORTANT:**
 - In boosting there are two different types of weights.
 - An *alpha* for each weak classifier that is chosen.
 - A weight for each training weight at each round.

Updating Training Weights

- So far we have chosen α_1 and h_1 .
- Next step: update training weights.
- Initially, training weights are uniform:
 - each weight is $1/(\text{number of training examples})$.
- After choosing h_1 , α_1 , each training weight is updated:

```
weights(i) = weights(i) * exp(-alpha * h1(i) * labels(i));
```

- When $\text{sign}(\alpha) * h(i) == \text{labels}(i)$:
 - we multiply weight by a number < 1 .

Updating Training Weights

- So far we have chosen α_1 and h_1 .
- Next step: update training weights.
- Initially, training weights are uniform:
 - each weight is $1/(\text{number of training examples})$.
- After choosing h_1 , α_1 , each training weight is updated:

```
weights(i) = weights(i) * exp(-alpha * h1(i) * labels(i));
```

- When $\text{sign}(\alpha) * h(i) \neq \text{labels}(i)$:
 - we multiply weight by a number > 1 .

Updating Training Weights

- So far we have chosen α_1 and h_1 .
- Next step: update training weights.
- Initially, training weights are uniform:
 - each weight is $1/(\text{number of training examples})$.
- After choosing h_1 , α_1 , each training weight is updated:

```
weights(i) = weights(i) * exp(-alpha * h1(i) * labels(i));
```

- Effect: examples that got misclassified by $\alpha_1 * h_1$ become more important.

Updating Training Weights

- So far we have chosen *alpha_1* and *h1*.
- Next step: update training weights.
- Initially, training weights are uniform:
 - each weight is $1/(\text{number of training examples})$.
- After choosing *h1*, *alpha_1*, each training weight is updated:

```
weights(i) = weights(i) * exp(-alpha * h1(i) * labels(i));
```

- Important: weights are always normalized to sum up to 1.
 - $\text{weights} = \text{weights} / \text{sum}(\text{weights})$.

Choosing the Second Weak Classifier

- Use function `find_best_classifier`.

```
find_best_classifier(responses, labels, weights)
```

- Choose an *alpha_2* for that classifier.

- Lower error \rightarrow higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- This is *exactly* the same thing that we did for choosing the first weak classifier.
- Why would we get a different weak classifier?

Choosing the Second Weak Classifier

- Use function `find_best_classifier`.

```
find_best_classifier(responses, labels, weights)
```

- Choose an *alpha_2* for that classifier.

- Lower error → higher weight.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- This is *exactly* the same thing that we did for choosing the first weak classifier.
- Why would we get a different weak classifier?
 - ***Because the training weights have changed.***

A Training Round of AdaBoost

- Try all classifiers.
- See which one has lower error rate.
 - IMPORTANT: error rates are measured on *weighted* training set.
- Choose an alpha for that classifier.
 - Lower error \rightarrow higher alpha.
- Update weights of training examples.

AdaBoost

- A simple for-loop, performing training rounds one at a time.

```

function result = AdaBoost(responses, labels, rounds)

result = zeros(rounds, 3);
classifier_number = size(responses, 1);
example_number = size(responses, 2);
weights = ones(example_number, 1) / example_number;
boosted_responses = zeros(example_number, 1);

for round = 1:rounds
    % find index, threshold, and alpha of best classifier
    [best_classifier, best_error, threshold, alpha] = ...
        find_best_classifier(responses, labels, weights);
    result(round, 1:3) = [best_classifier, alpha, threshold];

    % get outputs of the weak classifier on training examples
    weak_responses = double([responses(best_classifier, :)]' > threshold);
    weak_responses(weak_responses == 0) = -1;

    % reweigh training objects;
    for i = 1:example_number
        weights(i) = weights(i) * exp(-alpha * weak_responses(i) * labels(i));
    end
    weights = weights / sum(weights);

    % update boosted responses
    boosted_responses = boosted_responses + alpha * weak_responses;
    thresholded = double(boosted_responses > 0);
    thresholded(thresholded == 0) = -1;
    error = mean(thresholded ~= labels);
    disp([round error best_error best_classifier alpha threshold]);
end

```

Key Intuition

- Objects that are misclassified by the first chosen weak classifiers become more important.
- The next weak classifiers try to correct the mistakes of the first weak classifiers.

A Note on AdaBoost formulas.

- The actual formulas for AdaBoost may seem ad hoc:
 - Formula for choosing classifier alpha.

```
alpha = direction * 0.5 * log((1 - error) / error);
```

- Formula for updating training weights.

```
weights(i) = weights(i) * exp(-alpha * h1(i) * labels(i));
```

- These formulas can be justified mathematically.
 - That is beyond the scope of our class.


Applying Boosted Classifier

- Given an image:
- For every scale:
 - Compute integral image at that scale.
 - For every window:
 - For every weak classifier:
 - Compute response of window on that weak classifier.
 - Sum up the responses.
- For every pixel, get the best result among all windows (at all scales) centered at that pixel.
- See function `boosted_detector_demo`

Why Is it Useful

- We don't need to break our heads over:
 - What features to use.
 - How to combine them.

Results



How to Improve Accuracy

- More classifiers.
- More training data.
- Multiple training sessions.
 - After each session, find non-face windows that are detected as faces, and include to the training set.

How to Improve Efficiency

- Use a *cascade of classifiers*.
- First, try a strong classifier containing only few weak classifiers.
 - A lot of non-face windows can be rejected with a suitable threshold.
- What remains, is evaluated by a more complicated strong classifier.
- CLASS PROJECT