

Canny Edge Detection

CSE 6367 – Computer Vision
Vassilis Athitsos
University of Texas at Arlington

What Is an Edge?

- An edge pixel is a pixel at a “boundary”.
- There is no single definition for what is a “boundary”.



input



output1



output2

Image Derivatives

- We have all learned how to calculate derivatives of functions from the real line to the real line.
 - Input: a real number
 - Output: a real number
- In taking image derivatives, we must have in mind two important differences from the “typical” derivative scenario:
 - Input: two numbers (x, y) , not one number.
 - Input: discrete (integer pixel coordinates).
 - Output: Integer between 0 and 255.

Directional Derivative

- Let $f(x, y)$ be a function mapping two *real* numbers to a real number.
- Let θ be a direction (specified as an angle from the x axis).
- Let (x_1, y_1) be a specific point on the plane.
- Define $g(x) = f(x_1 + x \cos(\theta), y_1 + x \sin(\theta))$.
- Then, $g(x)$ is a function from the real line to the real line.
- The directional derivative of f at x_1, y_1 is defined to be $g'(0)$.

dx, dy Directional Derivatives

- For the directional derivative of f along the x axis, we use notation df/dx .
- For the directional derivative of f along the y axis, we use notation df/dy .

Vertical and Horizontal Edges

- Consider the image as a function $f(i,j)$ mapping pixels to intensity values.
 - Function $f(i,j)$ can be seen as a discretized version of a more general function $g(y,x)$, mapping pairs of real numbers to intensity values.
 - Vertical edges correspond to points in g with high dg/dx .
 - Horizontal edges correspond to points in g with high dg/dy .

Approximating dg/dx via Filtering

- In the discrete domain of $f(i,j)$, dg/dx is approximated by filtering with the right kernel:

```
dx = [-1 0 1;  
      -2 0 2;  
      -1 0 1];  
dx = dx / (sum(abs(dx(:))));  
dxgray = abs(imfilter2(gray, dx));
```

- Interpreting `imfilter(gray, dx)`:
 - Results far from zero (positive and negative) correspond to strong vertical edges.
 - These are mapped to high positive values by `abs`.
 - Results close to zero correspond to weak vertical edges, or no edges whatsoever.

Result: Vertical/Horizontal Edges

```
gray = read_gray('data/hand20.bmp');  
dx = [-1 0 1;  
      -2 0 2;  
      -1 0 1];  
dx = dx / (sum(abs(dx(:))));  
dy = dx'; % dy is the transpose of dx  
dxgray = abs(imfilter(gray, dx, 'symmetric', 'same'));  
dygray = abs(imfilter(gray, dy, 'symmetric', 'same'));
```



gray



dxgray
(vertical edges)



dygray
(horizontal edges)

Blurring and Filtering

- To suppress edges corresponding to small-scale objects/textures, we should first blur.

```
% generate two blurred versions of the image, see how it
% looks when we apply dx to those blurred versions.
filename = 'data/hand20.bmp';
gray = read_gray(filename);
dx = [-1 0 1; -2 0 2; -1 0 1] / 8;
dy = dx';

blur_window1 = fspecial('gaussian', 19, 3.0); % std = 3
blur_window2 = fspecial('gaussian', 37, 6.0); % std = 6
blurred_gray1 = imfilter(gray, blur_window1, 'symmetric');
blurred_gray2 = imfilter(gray, blur_window2, 'symmetric');
dxgray = abs(imfilter(gray, dx, 'symmetric'));
dxb1gray = abs(imfilter(blurred_gray1, dx, 'symmetric'));
dxb2gray = abs(imfilter(blurred_gray2, dx, 'symmetric'));
```

Blurring and Filtering: Results

gray



- Smaller details are suppressed, but the edges are too thick.
 - Will be remedied in a few slides, with *non-maxima suppression*.

dxgray

No blurring




dxblgray

Blurring, std = 3



dxblgray

Blurring, std = 6



Finding Edges at Other Angles

- Extracting edges at angle θ :
 - Rotate dx by θ , or
 - Rotate image by $-\theta$.
 - Rotating filter is typically more efficient.

`% detecting edges with orientation 45 or 135 degrees:`

```
fcircle = read_gray('data/blurred_fcircle.bmp');  
dx = [-1 0 1; -2 0 2; -1 0 1] / 8;  
rot45 = imrotate(dx, 45, 'bilinear', 'loose');  
rot135 = imrotate(dx, 135, 'bilinear', 'loose');  
  
edges45 = abs(imfilter(fcircle, rot45, 'symmetric'));  
edges135 = abs(imfilter(fcircle, rot135, 'symmetric'));
```

Results: Edges at Degrees 45/135

fcircle



edges45



edges135

More Edges at 45/135 Degrees

original image



edges at 45 degrees



edges at 135 degrees

Computing Gradient Norms

- Let:
 - $dxA = \text{imfilter}(A, dx);$
 - $dyA = \text{imfilter}(A, dy);$
- Gradient norm at pixel (i,j):
 - The norm of vector $(dxA(i,j), dyA(i,j))$.
 - $\text{sqrt}(dxA(i,j)^2 + dyA(i,j)^2)$.
- The gradient norm operation identifies pixels at all orientations.
- Also useful for identifying smooth/rough textures.

Computing Gradient Norms: Code

```
gray = read_gray('data/hand20.bmp');  
dx = [-1 0 1; -2 0 2; -1 0 1] / 8;  
dy = dx';  
  
blurred_gray = blur_image(gray, 1.4 1.4);  
dxgray = imfilter(blurred_gray, dx, 'symmetric');  
dygray = imfilter(blurred_gray, dy, 'symmetric');  
  
% computing gradient norms  
grad_norms = (dxblgray.^2 + dyblgray.^2).^0.5;
```

- See following functions online:
 - gradient_norms
 - blur_image


Gradient Norms: Results



gray



grad_norms



dxgray



dygray

Notes on Gradient Norms

- Gradient norms detect edges at all orientations.
- However, gradient norms in themselves are not a good output for an edge detector:
 - We need thinner edges.
 - We need to decide which pixels are edge pixels.

Non-Maxima Suppression

- Goal: produce thinner edges.
- Idea: for every pixel, decide if it is maximum along the direction of fastest change.
 - Preview of results:



gradient norms



result of nonmaxima suppression

Nonmaxima Suppression Example

- Example:

```
img = [112  118  111  115  112  115  112
       120  124  128  128  126  128  126
       132  134  132  130  130  130  130
       167  165  163  162  161  162  161
       190  192  199  196  198  196  198
       203  205  203  205  207  205  207
       212  214  216  219  216  213  217];
```



```
grad_norms = round(gradient_norms(img));
grad_norms = [ 5   5   7   7   7   7   7
               10  9   9   9   8   8   9
               23 21  18  17  17  17  17
               29 30  32  33  34  34  34
               19 20  20  22  22  22  23
               11 11  10  10  10   9   9
               5   5   6   6   5   4   5];
```



Nonmaxima Suppression Example

- Should we keep pixel (3,3)?
- result of dx filter $[-0.5 \ 0 \ 0.5]$
 - $(\text{img}(3,4) - \text{img}(3, 2)) / 2 = -2.$
- result of dy filter $[-0.5; 0; 0.5]$
 - $(\text{img}(4,3) - \text{img}(2, 3)) / 2 = 17.5.$
- Gradient = $(-2, 17.5).$
- Gradient direction:
 - $\text{atan2}(17.5, -2) = 1.68 \text{ rad} = 96.5 \text{ deg}.$
 - Unit vector at gradient direction:
 - $[0.9935, -0.1135]$ (y direction, x direction)



Nonmaxima Suppression Example

- Should we keep pixel (3,3)?
- Gradient direction: 96.5 degrees
 - Unit vector: $\text{disp} = [0.9935, -0.1135]$.
 - disp defines the direction along which pixel(3,3) must be a local maximum.
 - Positions of interest:
 - $[3,3] + \text{disp}$, $[3,3] - \text{disp}$.
 - We compare $\text{grad_norms}(3,3)$ with:
 - $\text{grad_norms}(3.9935, 2.8865)$, and
 - $\text{grad_norms}(2.0065, 3.1135)$



Nonmaxima Suppression Example

- We compare $\text{grad_norms}(3,3)$ with:
 - $\text{grad_norms}(3.9935, 2.8865)$, and
 - $\text{grad_norms}(2.0065, 3.1135)$
- $\text{grad_norms}(3.9935, 2.8865) = ?$
 - Use bilinear interpolation.
 - $(3.9935, 2.8865)$ is surrounded by:
 - $(3,2)$ at the top and left.
 - $(3,3)$ at the top and right.
 - $(4,2)$ at the bottom and left.
 - $(4,3)$ at the bottom and right.



Nonmaxima Suppression Example

- $\text{grad_norms}(3.9935, 2.8865) = ?$
 - Weighted average of surrounding pixels.

```
top_left = image(top, left);
top_right = image(top, right);
bottom_left = image(bottom, left);
bottom_right = image(bottom, right);
wy = 3.9935 - 3 = 0.9935;
wx = 2.8865 - 2 = 0.8865;
result = (1 - wx) * (1 - wy) * top_left +
          wx * (1 - wy) * top_right +
          (1 - wx) * y * bottom_left +
          x * y * bottom_right;
```

- See function `bilinear_interpolation` online.

img



grad_norms



Nonmaxima Suppression Example

- $\text{grad_norms}(3.9935, 2.8865) = 33.3$
- $\text{grad_norms}(2.0065, 3.1135) = 10.7$
- $\text{grad_norms}(3,3) = 18$
 - Position 3,3 is not a local maximum in the direction of the gradient.
 - Position 3,3 is set to zero in the result of non-maxima suppression
 - Same test applied to all pixels.



Nonmaxima Suppression Result


grad_norms =

```
[ 5   5   7   7   7   7   7
 10   9   9   9   8   8   9
 23  21  18  17  17  17  17
 29  30  32  33  34  34  34
 19  20  20  22  22  22  23
 11  11  10  10  10   9   9
 5   5   6   6   5   4   5];
```

img



grad_norms



nonmaxima_suppression(grad_norms, thetas, 1) =

```
[ 0   0   0   0   0   0   0
 0   0   0   0   0   0   0
 0   0   0   0   0   0   0
 0  29  34  33  34  33   0
 0   0   0   0   0   0   0
 0   0   0   0   0   0   0
 0   0   0   0   0   0   0]
```



result of
non-maxima
suppression

Nonmaxima Suppression Result



gradient norms



result of nonmaxima suppression

Side Note: Bilinear Interpolation

- `grad_norms(3.9935, 2.8865) = ?`
 - Weighted average of surrounding pixels.

```
top_left = image(top, left);
top_right = image(top, right);
bottom_left = image(bottom, left);
bottom_right = image(bottom, right);
wy = 3.9935 - 3 = 0.9935;
wx = 2.8865 - 2 = 0.8865;
result = (1 - wx) * (1 - wy) * top_left +
          wx * (1 - wy) * top_right +
          (1 - wx) * y * bottom_left +
          wx * wy * bottom_right;
```

- Interpolation is a very common operation.
 - Images are discrete, sometimes it is convenient to treat them as continuous values.

bilinear_interpolation.m

```
function result = bilinear_interpolation(image, row, col)

% row and col are non-integer coordinates, and this function
% computes the value at those coordinates using bilinear interpolation.

% Get the bounding square.
top = floor(row);
left = floor(col);
bottom = top + 1;
right = left + 1;

% Get values at the corners of the square
top_left = image(top, left);
top_right = image(top, right);
bottom_left = image(bottom, left);
bottom_right = image(bottom, right);

x = col - left;
y = row - top;

result = (1 - x) * (1 - y) * top_left;
result = result + x * (1 - y) * top_right;
result = result + x * y * bottom_right;
result = result + (1 - x) * y * bottom_left;
```

The Need for Thresholding




gray

- Many non-zero pixels in the result of nonmaxima suppression represent very weak edges.



nonmaxima



nonmaxima > 0

The Need for Thresholding




gray

- Decide which are the edge pixels:
 - Reject maxima with very small values.
 - Hysteresis thresholding.



gradient norms



result of nonmaxima suppression

Hysteresis Thresholding

- Use two thresholds, t_1 and t_2 .
- Pixels above t_2 survive.
- Pixels below t_1 do not survive.
- Pixels $\geq t_1$ and $< t_2$ survive if:
 - They are connected to a pixel $\geq t_2$ via an 8-connected path of other pixels $\geq t_1$.

Hysteresis Thresholding Example



A = nonmaxima ≥ 4



B = nonmaxima ≥ 8



C = hystresh(nonmaxima, 4, 8)

- A pixel is white in C if:
 - It is white in A, and
 - It is connected to a white pixel of B via an 8-connected path of white pixels in A.

Canny Edge Detection

- Blur input image.
- Compute dx, dy, gradient norms.
- Do non-maxima suppression on gradient norms.
- Apply hysteresis thresholding to the result of non-maxima suppression.
 - Check out these functions online:
 - blur_image
 - gradient_norms
 - gradient_orientations
 - nonmaxsup
 - hysthresh
 - canny
 - canny4

Side Note:

Angles/Directions/Orientations

- To avoid confusion, you must specify:
 - Unit (degrees, or radians).
 - Do you use undirected or directed orientation?
 - Undirected: $180 \text{ degrees} = 0 \text{ degrees}$.
 - Directed: $180 \text{ degrees} \neq 0 \text{ degrees}$.
 - Which axis is direction 0? Pointing which way?
 - Class convention: direction 0 is x axis, pointing right.
 - Do angles increase clockwise or counterclockwise?
 - Class convention: clockwise.
 - Does the y axis point down? (in this class: yes)
 - What range of values do you allow/expect?
 - $[-180, 180]$? Any real number?

Side Note:

Angles/Directions/Orientations

- Confusion and bugs stemming from different conventions are extremely common.
- When combining different code, **make sure that you account for different conventions.**

```
thetas = atan2(dyblgray, dxb1gray);
```

```
% atan2 convention:  
% values: [-pi, pi]  
% 0 degrees: x axis, pointing left  
% y axis points down  
% values increase clockwise
```

Side Note: Edge Orientation

- How is the orientation of an edge pixel defined?

Side Note: Edge Orientation

- How is the orientation of an edge pixel defined?
 - It is the direction PERPENDICULAR to the gradient, i.e., the (dx, dy) vector for that pixel.
 - Typically (not always) 0 degrees = 180 degrees.
 - In other words, typically we do not care about the direction of the orientation.